Price Incentives for Conservation:
Experimental Evidence from Groundwater Irrigation

Nick Hagerty

Ariel Zucker [*]

June 28, 2025

Abstract

Policymakers often seek to reduce resource consumption but face constraints that preclude the use of prices or other first-best instruments. This paper studies an alternative approach: payments for voluntary conservation. We evaluate a randomized controlled trial of payments for reduced groundwater use among farmers in Gujarat, India, where both water and the electricity used to pump it are scarce and largely unregulated. Price incentives work: The program reduced hours of irrigation by 22 percent and energy use by 23 percent. The cost of payments per unit energy conserved was comparable to the cost of electricity supply, suggesting that conservation payments are a practical and cost-effective policy tool for electric utilities. A critical program design feature is the benchmark used to measure conservation: more stringent benchmarks reduced inframarginal payments, induced greater conservation among high users, and halved the average cost of conservation.

1 Introduction

Across a wide range of settings, natural resources – and the externalities associated with their use – are unpriced. Policymakers and other stakeholders often seek to reduce resource consumption but face constraints that rule out first-best instruments. Pigouvian taxes may be politically untenable, while creating new property rights or markets may be practically infeasible. Common second-best alternatives, such as rationing or subsidies for resource-saving technologies, tend to be inefficient or cost-ineffective. There is a need for alternative tools that can achieve conservation goals while operating within real-world political and administrative constraints.

One alternative that is gaining wider use is payments for voluntary conservation, in which a government (or other principal) offers payments to users who reduce their resource consumption. Such a program can replicate the marginal incentives of a Pigouvian tax, at least for some users, so it may achieve similar outcomes. But because it relies on rewards rather than penalties, it reverses the distributional consequences, potentially relaxing policy constraints.

This paper experimentally evaluates a program of payments for groundwater conservation among farmers in Gujarat, India – a setting where marginal incentives to conserve are minimal, yet the consequences of overuse are severe. Groundwater is an essential resource for irrigation and drinking water worldwide, but unregulated extraction in many regions has led to depletion, increasing poverty and conflict while reducing farm income, wealth, and employment (Sekhri, 2014; Blakeslee et al., 2020). The problem is compounded by electricity pricing: power used to pump groundwater is often free, flat-rate, or heavily subsidized, exacerbating depletion while straining public utilities and degrading electricity service (Burgess et al., 2020). Yet high energy subsidies are often seen as valuable means of redistribution; in India, reform efforts are often met with political resistance and public protest (Sovacool, 2017).

Our program installed meters and offered payments for reduced groundwater pumping between Fall 2022 and Spring 2023 within a randomized controlled trial. The basic design is to (1) meter the groundwater pumps of all study participants, (2) offer randomly selected participants payments for reduced pumping relative to a "benchmark" quantity, and (3) compare groundwater consumption by these farmers to that of the rest of the sample. The program was implemented in collaboration with the Aga Khan Rural Support Programme (India), a trusted organization with a long history in the study area.

We first ask how this program affected groundwater and electricity consumption. We find that payments for groundwater conservation work. Farmers offered conservation

payments irrigated for 22 percent fewer hours than farmers assigned to the control group, with a 95 percent confidence interval of $(12, 33)$ percent. The effects are similar when we convert irrigation time to energy use—we estimate that treatment farmers reduced electricity consumption by 140 kilowatt-hours (kWh) per month relative to the control-group mean of 611 kWh. Treatment effects increased over the three months of the intervention, suggesting a durable response. Higher prices (relative to lower prices) reduced electricity use but not pumping duration, suggesting that farmers with more powerful pumps are more price-sensitive.

Then, we assess the cost-effectiveness of payments for groundwater conservation from the perspective of an electric utility – the natural choice to implement a similar program at scale. At the margin, is it cheaper for a utility to increase supply or to reduce demand through conservation payments? We find that our program is cost-competitive with typical costs of electricity procurement in northwestern India. Our program spent 6.1 INR in total conservation payments for every kWh of energy saved, a cost that is slightly greater than the average costs of electricity provision for the utility in our study area (and may be lower than marginal costs) and slightly lower than costs in a nearby state. Considering the additional social costs of groundwater depletion and of emissions from electricity generation, it appears likely that paying farmers to reduce their groundwater irrigation would bring greater social benefits than purchasing more electricity and distributing it for free.

Finally, we study the optimal design of contracts in our setting. How should we set the benchmarks – the values against which conservation is measured and rewarded? Often a natural choice is the value of the conservation outcome prior to the program, since it may be a good predictor of what the outcome would later be absent the program. But what if the counterfactual cannot be predicted well? We argue that conservation programs should not try to set the benchmark accurately; they should instead set the benchmark cost-effectively. We also argue that cost-effective benchmarks are likely to be low, or deliberately ambitious.

The reason is that the right objective function is asymmetric and nonlinear in the prediction error. First, if benchmarks are too high, they generate excessive inframarginal payments – rewards for units of conservation that never would have been conserved anyway. We show in our setting that inframarginal payments are large, but would have been even larger had we not set low benchmarks and targeted them on expected variance. Second, low benchmarks can actually induce greater conservation. In our treatment group, we also cross-randomized participants into high- and low-benchmark groups, which separated benchmarks by about 30 percent. Many participants responded to lower bench-

3

marks by conserving more, and the average cost of conservation in the high-benchmark group was more than twice the average cost in the low-benchmark group.

One central contribution of this paper is to provide experimental evidence that marginal prices can reduce groundwater irrigation. The basic idea is core to microeconomic theory, but empirical evidence has been limited because there have been so few real-world examples to study. We build most directly upon two non-randomized pilots of similar programs run by electric utilities in India; Fishman et al. (2016) found no effects in Gujarat, while Mitra et al. (2023) found water use reductions in Punjab. By working instead with a local non-governmental organization, we are able to randomize participants, directly measure pumping at the individual level, and trace the demand curve to prices beyond what utilities have been willing to test. In the most closely related randomized study, Chakravorty et al. (2023) induce a volumetric price using an encouragement design; they do not find effects on water use, which may be because relatively few farmers adopted the price.[1]

Observational research on irrigation prices has faced two key challenges: finding exogenous variation in groundwater prices, and reliably measuring water use. Early studies in the U.S. used proxies for pumping costs and relied on self-reported water use (Gonzalez-Alvarez et al., 2006; Hendricks and Peterson, 2012; Pfeiffer and Lin, 2014). A few more recent papers study the introduction of explicit prices. One in West Bengal uses self-reports of water use and finds inconclusive results (Meenakshi et al., 2013). Three in the U.S. are able to observe meter readings from individual wells and find reductions in water use (Smith et al., 2017; Bruno and Jessoe, 2021; Bruno et al., 2024). Introducing randomized variation allows us to more cleanly isolate causal effects, while installing meters allows us to directly observe irrigation behavior in a setting beyond the western U.S.

Our project also contributes to the broader literature on payments for environmental services (PES). Our intervention has the same basic structure as hundreds of programs designed to incentivize the provision of environmental services, ranging from increased forest or wetland cover, to reduced input intensity in agriculture.[2] Despite their prevalence, rigorous evaluation of these types of programs has been limited (see Pattanayak et al. (2010) and Börner et al. (2017) for reviews). Most existing evaluations use covariate matching and are unable to address selection bias, a particular concern for a voluntary program. The exceptions are three randomized controlled trials of programs to reduce

---

[1]They also show that a water-saving technology saves water only in regions with marginal prices; although this technology is randomized, the heterogeneity is not.

[2]For example, in the United States alone, payments are available to farmers for actions to mitigate flood and wildfire risks, provide habitat for endangered species, salinity mitigation, and water and energy conservation.

deforestation in Uganda (Jayachandran et al., 2017), encourage tree planting in Malawi (Jack and Cardona Santos, 2017), and reduce crop burning in India (Jack et al., 2023). Our study shows that PES models are feasible and can be effective in a novel context: reducing energy and water use in agriculture.

Finally, we contribute to literature connecting the price response of electricity consumption in developing countries to policy decisions about energy-sector investment and reform. Experimental and quasi-experimental studies are still limited, but a few have been conducted recently on rural households in Columbia (McRae, 2015), urban households in South Africa (Jack and Smith, 2016), and new grid connections in Kenya (Lee et al., 2018).

2 Study Setting and Experimental Design

2.1 Enrollment and Sample

We implemented a randomized controlled trial among groundwater-irrigating farmers in Saurashtra, a water-scarce region of Gujarat state, India. The study villages are located in the inland districts of Rajkot, Surendranagar, and Morbi (shown in Figure 1). Groundwater depletion is a concern within the study area, and nearby areas are marked by some of the most rapid groundwater depletion rates both within India and globally (Jasechko et al., 2024). While the primary source of employment in the study area is in agriculture (Registrar General and Census Commissioner of India (2001)), there are also a number of industrial occupation opportunities.

We recruited our sample using lists of villagers currently or formerly participating in agricultural outreach programs with our implementing partner, the Aga Khan Rural Support Program (AKRSP), and its sister agency the Aga Khan Foundation (AKF). The outreach programs included Better Cotton Initiative, which aims to improve the sustainability of the global cotton supply; Farmer Producer Groups, which aim to empower farmers in marketing produce and procuring high-quality inputs; and various micro-irrigation subsidy and support programs. Surveyors approached farmers on these lists, as well as any farmers who shared water with those on the lists, to determine eligibility.

In order to be eligible for the study, the household's primary agricultural decision-maker (PAD) was required to meet the following criteria: Planted crops and irrigated primarily using groundwater in the previous Rabi season; planned to irrigate during the coming Rabi season; had no more than two active wells on their primary farm; had electric-powered pumps on all active wells; did not have multiple pump starters in use

5

on any active well; and did not belong to a network of sharing irrigation sources among groups of farmers larger than four.

We enrolled a total of 1,347 farmers who met the eligibility criteria, completed a baseline survey, and consented to the full study (including installation of an hours of use meter on the pumpsets used to irrigate their primary farm). Of these, 236 attrited prior to randomization, and another 122 prior to the final data collection visit, leaving an analysis sample of 989 farmers.

2.2   Interventions

The experiment had two overarching treatment arms: *conservation credit* farmers were eligible to receive payments for conserving groundwater below a benchmark, whereas *control* farmers received no such incentives.

Conservation credits   Our study utilizes random variation in whether participants were assigned to four versions of the Conservation Credits intervention arms, as shown in Figure 2. Participants in the Conservation Credits arms had an hours-of-use meter installed on the electric pump starter of their primary irrigation source or sources.[3] The meter measures the total hours of irrigation done by the farmer.

Meters were installed in Fall of 2022, and were read monthly by survey staff from December 2022 through March 2023 (Figure 3). We took three actions to discourage tampering and removal. First, stickers were placed at the easiest disconnection points such that disconnection would tear the sticker, enabling easy detection. Second, in the case that meter removal or tampering was detected, participants were disqualified from receiving conservation incentives. Finally, participants were rewarded 100 INR per meter for keeping their meters installed without tampering through the final meter reading.

Farmers in the Conservation Credits arms were informed of their treatment assignment at the first meter reading in December 2022. Farmers were incentivized for conserving water for the following four months of the winter growing season, known as the Rabi season, from December-March. This is the period of peak irrigation; as there is typically no rainfall during Rabi, agriculture is entirely dependent on irrigation. At each meter reading, farmers were informed of their benchmark for the following month, and the payment for the previous month was calculated. Payments were awarded at a fixed rate for consuming fewer hours of irrigation than the monthly benchmark, according the

---

[3]Farmers in our sample had up to two wells on their primary farm, and therefore up to two metered pump starters.

formula:

$$\text{Payment}_{it} = \min\left(\max\left(0, \text{price}_i \times \left(\left(\text{hours benchmark}\right)_{it} - \left(\text{hours consumed}\right)_{it}\right)\right), \left(\text{max payment}\right)_i\right) \quad (1)$$

where price_i is the per-hour incentive rate, $\left(\text{hours benchmark}\right)_{it}$ is an individual-month-specific benchmark, $\left(\text{hours consumed}\right)_{it}$ is the monthly meter reading, and $\left(\text{max payment}\right)_i$ is the maximum monthly payment.[4] Payments were pro-rated in the case that meter readings were not exactly 31 days apart. The payments were later disbursed via electronic bank transfer.

Conservation Credit Sub-treatments   The four Conservation Credits sub-treatments differ along two dimensions: the per-hour incentive rate, and the benchmark. Individuals assigned a *high price* received 100 INR (1.20 USD) per hour conserved, and those assigned a *low price* received 50 INR (0.60 USD) per hour conserved. The prices were chosen to encompass realistic ranges of groundwater prices that a policymaker might wish to set. The low price represents the approximate cost of electricity provision for the median farmer and is similar to the price offered in a program in Punjab (Mitra et al., 2023).[5] The high price allows us to study the response to prices well beyond those piloted by electric utilities in India to date, which might be justified by the additional social costs of groundwater depletion and electricity generation.

Individualized benchmarks were set using a formula that optimized the expected number of marginal farmers as a function of first-month pumping data (i.e., after meter installation but before treatment assignment was revealed to surveyors or farmers). Individuals assigned the *high* and *low benchmark* received 115% and 85%, respectively, of their formula-based benchmark, rounded to the nearest 10-hour increment.

Control   Participants in the Control arm also had hours-of-use meters installed and read monthly for four months, and were rewarded 100 INR per meter for keeping the meters installed for the duration of the intervention. However, these farmers were not incentivized for conservation.

---

[4]The maximum monthly payment was 4,000 INR for farmers with one well and 6,000 INR for farmers with two wells. These maximums were not pro-rated.

[5]A price of 50 INR per hour is approximately equal to the unsubsidized average cost of electricity supply in Gujarat for the median pumpset in our sample. That is: (5.4 INR/kWh average cost of electricity provision in Gujarat) * (5 hp pump brake power) / (40% typical motor efficiency) * (0.75 kW/HP conversion factor) = 50 INR/hr. The Punjab program offered an incentive of 4.0 INR/kWh, which translates to different per-hour prices for different farmers depending on pump power, but would be approximately 37 INR per hour for the median pumpset in our sample. For more details on this calculation see Section 3.2 and Table 5.

2.3 Randomization

Randomization was conducted at the level of farmer-sharing group: that is, the set of farmers who mentioned at the baseline survey that they used any common irrigation sources. By randomizing at the sharing group level, we minimize the possibility that conservation credits will spillover to control farmers.

Randomization was stratified by forecasted hours of irrigation and size of water-sharing group. Specifically, the final sample of water-sharing groups was ordered first by number of farmers, and second by forecasted hours of irrigation.[6] Groups were randomly allocated in equal proportion between the Control and Conservation Credit arms using a pseudo-random number generator (Stata software) within each ordered pair. Pairs were then combined into ordered cells of eight farmer-sharing groups, within which the four groups allocated the Conservation Credit arm were randomly allocated in equal proportion between the four Conservation Credits sub-treatments.

3 Data and Summary Statistics

3.1 Data Sources

Our analysis rests on data from two primary sources: a baseline survey and meter reading data. First, we conducted a baseline survey with both self-reported and field measurement components prior to randomizing participants into treatments. Self-reported data include demographic and socioeconomic characteristics, such as landholding size and household size; cropping, crop management, and irrigation decisions in the previous year; the power of the primary pumpsets; and water conservation strategies and attitudes. Field measurements include the precise geolocation and depth-to-water of each well on the participant's largest farm where measurement is safe and feasible. We also collect the names and contact details of any farmers who use water from the primary farm or whose water is used on the primary farm in order to sort our sample into water-sharing groups consisting of all farmers who are connected through water sharing relationships. Baseline data was collected electronically through tablet surveys.

Second, we directly measure groundwater pumping for all study participants using hours-of-use meters installed on the pump starter of each participant's primary irrigation source.[7] Surveyors recorded meter readings each month using a digital tablet survey.

---

[6]Forecasted hours of irrigation were created using a random forest using baseline survey data and geological mapping data. Forecasts were fit using a sample of farmers in Saurashtra from a previous project.

[7]Analog hours-of-use meters manufactured by Nishant Engineers (model: NE53/6S).

Meter data quality was assured through random audits, in which a research associate compared the digitally recorded meter readings with dated, geo-located photographs of the meter dial included on the tablet survey.

We combine these with two secondary data sources: hydrogeology data from digitized "Groundwater Prospects Maps" created by the Government of India and proxies of agricultural yields derived from satellite data. Specifically, we first supplement our baseline data with an extensive set of hydro-geological features from the "Groundwater Prospects Maps" prepared by the National Remote Sensing Center, Government of India. These maps are available only as images; we digitize the map images and convert them to vector data that can be used in analysis. Following Ryan and Sudarshan (2022), we extract variables on rock type, aquifer type, and fractures that are predictive of groundwater availability and yields in hard rock areas aquifer systems.

We supplement both our baseline and outcome data using agricultural yield proxies derived from satellite imagery. We utilize two common proxies for agricultural biomass: the Enhanced Vegetation Index (EVI) and Normalized Difference Vegetation Index (NDVI). These indices are calculated using Google Earth Engine from Sentinel-2 images at a 10 meter resolution (which are taken every five days) in a 40 meter radius around the coordinates of participants' wells. We filter out pixels of clouds and surface water, and use standard techniques to correct for data anomalies from atmospheric disturbances.

3.2 Outcome Variables

Our primary outcome variable is monthly hours of groundwater irrigation. Meters show cumulative duration of pump operation, so we calculate monthly irrigation hours as the difference between values shown on the meter at the current visit and previous month's visit. Because not all meter-reading visits occurred at exact monthly intervals, we rescale observed hours to a 31-day rate so that observations are comparable across farmers and months.

Our secondary outcome variable is energy consumption in irrigation. Energy use is not observed directly but rather converted from hours of irrigation using known functional relationships from physics. The formula is:

$$E = \frac{P_b}{\eta_m} \times t \tag{2}$$

where E is energy consumed, t is duration of pump operation, P_b is the power rating

of the pump's motor ("brake horsepower"), and η_m is the motor efficiency, a unitless constant between zero and one.[8]

We collect t and P_b in meter-reading and baseline surveys. Motor efficiency η_m is difficult to measure accurately and would have required use of electricity meters, with which many study participants were uncomfortable. Instead, we draw an estimate of motor efficiency from the literature in the most similar setting we can find: 40 percent (Mitra et al., 2023).

Note that our energy use variable is not simply a monotone transformation of irrigation hours, since it also depends on pump power, which varies across farmers. That said, pump power does not endogenously respond to the program,[9] so energy use can be seen as a rescaling of units combined with a reweighting of farmers within the sample. Either way, average treatment effects may be substantively different if individual treatment effects are heterogeneous and correlated with pump power. Similarly, individual-level price elasticities of demand for energy and hours are equal, but aggregate price elasticities may differ due to this reweighting.

A third set of outcomes measures the economic impacts of water conservation using two proxies for agricultural yields: log seasonal differences of EVI and NDVI for each farmer. For each index, we follow Asher and Novosad (2020) and proxy productivity over the 2022-23 rabi season (December – March) as the maximum index value reached on the farmers' land during the season less mean index value in the first four weeks of the season.[10] Our preferred outcome is the natural log of this difference, which eases interpretation. We calculate rabi productivity using the same methodology in the year prior to the intervention (2021-22) and mean index values in November 2022, prior to randomization, as additional baseline controls.

3.3 Descriptive Statistics and Balance Checks

Descriptive statistics. Table 1 reports baseline characteristics of the experimental sample. In both this table and all subsequent analysis, the sample is restricted to farmers who completed all rounds of data collection: the baseline survey, meter installation, the

---

[8]To obtain E in kilowatt-hours (kWh) when P_b is measured in horsepower (hp), the formula also requires the unit conversion of 0.7457 kW per hp.

[9]In a longer-term incentive program, farmers could make investments in new pumps. Our program lasted for only one irrigation season, and participants were unable to replace a pump without removing the meter and becoming disqualified from the program.

[10]Daily index values are first averaged across wells for farmers with multiple wells. We difference the mean index value of the first four weeks of the growing season, as opposed to six weeks in Asher and Novosad (2020), to avoid contamination by the start of our intervention.

baseline meter reading, and all three meter-reading visits during the intervention.

Our sample consists predominantly of smallholder farmers; the mean plot area of their primary farm is 1.95 hectares.[11] Most participants are literate, have completed primary and secondary education, and identify with a "scheduled caste/scheduled tribe/other backward caste" designation. Only half own a plow or tractor. Cotton is the primary crop in our sample, with sorghum/millet, groundnut, and pulses as the next most common crops. Farmers are at least somewhat diversified in their crops, with a mean count of distinct crops around 2.

Most participants have only one active well; some have two. Some wells are dug-wells and others are borewells (tubewells); the most common type of well is a dug-*cum*-borewell, in which a borehole is drilled into the bottom or side of a dugwell in order to access additional pockets of water. The average well is 59 meters deep, but many are considerably deeper. The most common electric pump installed in each well is rated at 5 horsepower, but some are more powerful.

Many farmers in our sample are already using cultivation practices that conserve water. 41 percent use a drip irrigation system, 69 percent use raised beds, and 19 percent use rotational, strip, or zero tillage. Local water markets are rare in our context: Only 1 percent report having purchased water for irrigation. Farmers sometimes share irrigation sources with neighbors, usually relatives, but water sharing is not a large share of irrigation in our sample: About 10 percent of pump operation during the previous (2021-22) irrigation season was directed to irrigation off the primary farm, which includes both neighbors and secondary farms also held by the respondent.

Balance.   Columns 3 and 4 of Table 1 report means of baseline characteristics separately for the overall treatment and control groups. The two groups appear similar across all characteristics. We formally check for balance test between the main treatment and control groups using a Wald *F*-test for joint orthogonality of all characteristics reported in this table. The *F*-statistic is small and the *p*-value is large, so we fail to reject the null hypothesis that treatment-control differences are zero for all characteristics.

Our sample includes more farmers in the treatment group than in the control group, implying that attrition rates were different across the two groups. Differential attrition would bias the results if attrition is correlated with characteristics that predict the outcome variable. But we do not see evidence that the treatment and control groups are differentially selected across a range of baseline characteristics.

---

[11]The Indian government typically defines farmers holding less than 2 hectares as "small and marginal farmers".

4 Program evaluation

We first evaluate our conservation payments program as implemented by estimating intent-to-treat (ITT) effects of eligibility for the overall intervention. We estimate post-double-selection LASSO regressions of the following form:

$$Y_{it} = \alpha_t + \tau \cdot \text{ConservationCredits}_i + \gamma' \mathbf{X}_{it} + \varepsilon_{it}, \tag{3}$$

where Y_{it} is a outcome variable for farmer i at monthly visit t, and $\text{ConservationCredits}_i$ is an indicator for being in the overall treatment group and therefore eligible for payments. X_{it} is a vector of individual-specific covariates interacted with month t chosen by double-selection LASSO (Belloni et al., 2013). [12] Standard errors are clustered by randomization pair (following de Chaisemartin and Ramirez-Cuellar (2024)), which nests months within farmer and farmers within groups of neighbors that reported sharing water prior to the intervention.[13]

Since our primary outcome variables are right-skewed, we also estimate the same specifications as Poisson regressions via pseudo-maximum likelihood, following Chen and Roth (2024):

$$Y_{it} = \exp\{\alpha_t + \tau \cdot \text{ConservationCredits}_i + \gamma' \mathbf{X}_{it}\} \cdot u_{it}, \tag{4}$$

Poisson regressions can more precisely estimate regression-adjusted treatment effects when covariates have multiplicative effects on the outcome. They directly recover a transformation of the average treatment effect as a proportion of the control mean (Silva and Tenreyro, 2006).[14]

4.1 Conservation payments reduce irrigation time and energy use

Table 2 presents estimated effects of overall payment eligibility on hours of irrigation, the variable directly measured by our meters. Our preferred estimate is the covariate-adjusted specification in column (3): Farmers assigned to the program operated their

---

[12]We implement post double-selection LASSO in Stata with the commands dsregress for OLS, dspoisson for Poisson, and the user-written command ivlasso (Ahrens et al., 2018) for instrumental variables regressions, respectively.

[13]We omit randomization pair fixed effects following Bai et al. (2023), who show that they complicate the interpretation of the estimand and do not necessarily reduce bias from differential attrition.

[14]Our primary outcome variables sometimes take a value of zero, so we cannot run log-linear OLS regressions. We avoid other "log-like" transformations, such as the inverse hyperbolic sine, because they are sensitive to the choice of units (Mullahy and Norton, 2023) and because any notion of an individual-level percentage change is undefined for a variable that admits zero (Chen and Roth, 2024).

pump for an average of 10.4 fewer hours per month during the intervention period than control-group farmers. This effect represents a 22 percent reduction relative to the control-group mean of 47 hours per month, and the 99% confidence interval excludes zero.

Results are broadly robust to alternative specifications. Column (7) shows the same specification as column (3) but estimated using Poisson regression, again indicating that the program led to a 22 percent reduction ($e^{-0.25} - 1$) in irrigation hours. Estimates are less precise without covariate adjustment (columns (1) and (5)) but the 90% confidence intervals still exclude zero. Columns (2) and (5) include village fixed effects and no other covariates, a parsimonious example of a fully saturated regression model. We include this specification to confirm that the main result is similar in a specification guaranteed to be unbiased for the average treatment effect even in finite samples (Athey and Imbens, 2017).

To see how the program affected the distribution of pumping hours, Figure 4 plots bin treatment effects – i.e., effects of payment eligibility on the share of participants whose pumping hours fall within certain ranges – using a subset of controls. These estimates are less precise but important patterns are visible. Figure 4(a) plots bin effects by hours of irrigation. It shows that the program moved participants from higher values of pumping to lower values. Figure 4(b) plots bin effects by hours relative to the benchmark, which varies across participants. It shows that the program moved participants from above their benchmark to below it, suggesting that participants understood the program structure and complied with its incentives.[15]

Effects on energy use are shown in Table 3. Again our preferred estimate is in column (3): Payment eligibility reduced energy use by 140 kWh per month, a 23 percent reduction compared with a control-group mean of 611 kWh per month. Alternative specifications are broadly consistent though less precise than the irrigation hours regressions. Point estimates without covariate adjustment are much smaller but so imprecise that we cannot reject equality with our preferred specification.

4.2 Treatment effects increased over time

To investigate seasonal patterns in treatment effects, we augment our primary regressions to estimate separate treatment effects in each month of the program. Results are plotted in Figure 5 for both OLS and Poisson estimates of our preferred specification, which includes

---

[15]Usually we would not expect to see bunching immediately below the benchmark threshold since the budget set is concave at this point; anti-bunching is more likely. In this case, the reason that effects concentrate in the first bin below the threshold is likely due to the zero lower bound. Most benchmarks are set in the 10-50 hour range, so most participants can never appear in the first three bins shown in the graph.

controls selected by double LASSO.

Average effects of the program increased in magnitude over the course of the experiment, from 5 hours in the first month to 12 and 13 hours in the last two months of the program. (We can reject that the first and second, or the first and third months, are equal, at a 5 percent significance level.) These differences are even more dramatic when expressed as a percentage of the control mean, which declined over time. Treatment effects estimated using Poisson regression increased from 9 percent in the first month to 22 and 28 percent in the second and third months.

We see two likely explanations for the growing response over time. One possible reason is increasing trust in the program. Because the conservation credits program was a completely new concept, it seems likely that participants would have changed their behavior only tentatively in the first month. After they saw real cash appear in their bank accounts, they responded less cautiously. Another possible reason is that demand for irrigation becomes more elastic later in the growing season. For many crops, water application is most critical during an early phase of growth. After this early phase, yields may be less sensitive to irrigation amounts, and so farmers would become more sensitive to the price of irrigation. We do not currently have data to distinguish between these explanations, but we expect both are at play.

4.3 Higher prices have little additional effect

Next, we go beyond the effects of the program overall to investigate whether the level of price incentive affects irrigation behavior, conditional on program participation. We compare the high- and low-price sub-treatment groups by interacting the overall treatment variable with an indicator for being in the high-price subgroup.The results are in columns (4) and (8) of Tables 2 and 3.

Across specifications and outcomes, the main effects of the program remain large and statistically significant, while the interaction effects are smaller and not statistically significant. This says that being offered a price incentive of 50 INR per hour, relative to not being offered a price incentive at all, has a greater effect on conservation than increasing the price from 50 to 100 INR per hour. This result is consistent with a convex demand curve: There may be many low-cost opportunities to conserve water and energy resources that are left on the table when marginal resource prices are zero but adopted when prices are positive, but once that low-hanging fruit is picked, resource conservation faces more rapidly rising opportunity costs.

While none of the interaction effects are statistically significant at the 5% level, the

magnitude of the point estimates vary considerably across specifications and outcomes. The Poisson estimate for hours of irrigation suggests that higher prices have no substantial independent effect, while the OLS estimate for the impact on higher prices on energy use (-111.2 kWh) is about two-thirds as large as the main effect of the program (-139.8). One explanation for this is that farmers with higher-powered pumps tend to respond more to higher prices than low. We can speculate that for these farmers, the opportunity cost of an hour of pumping is larger, and a higher hourly price may thus be needed to encourage conservation.

4.4 Conservation payments do not harm agricultural productivity

To understand whether the intervention has any measurable economic impacts, we use Equation 3 to examine the effects of the intervention on proxies of agricultural yields. Results are shown in Appendix Table 11.

Despite reducing irrigation intensity, payment eligibility does not lead to any significant change in yields. Our preferred estimate is in column (3): the seasonal increase in NDVI among payment eligible farmers is about 7.6 percent *larger* than among control farmers, and although this estimate is noisy we can reject substantial decreases in productivity among treated farmers. The results are similar if we use alternative measures of agricultural productivity. Overall, we find no evidence that the reductions in irrigation had negative economic consequences for farmers.

5 Demand estimation

We now use the experimental variation introduced by our program to estimate the slope of demand for groundwater irrigation. The idea is that in a program of payments for voluntary conservation, not all farmers are actually marginal to the incentive, unlike as they would be under a universal volumetric electricity price or groundwater pumping fee. Even for farmers offered payments, the marginal price is zero for those who pump for more hours than the benchmark, as well as for those who reach the maximum payment. Figure 6 illustrates these cases relative to the budget set created by a conservation payments program.

As a result, the treatment effect depends on specific design parameters of our program: price, benchmarks, and maximum payments. In contrast, a demand model gives us potentially more generalizable information as to how the farmers in our sample would adjust their irrigation behavior under other types of programs.

To estimate demand, we estimate instrumental variables regressions of irrigation on price, instrumenting for price with the experimental treatment groups:

$$Y_{it} = \alpha_t + \beta p_{it} + \gamma' \mathbf{X}_{it} + \varepsilon_{it}, \tag{5}$$

where $p_{it} \in \{0, 50, 100\}$ represents the effective marginal price of an hour of irrigation faced by farmer i in month t. Effective marginal price in each month is zero for control-group farmers, for treatment-group farmers who did not receive a payment, and for treatment-group farmers who reached the maximum payment. For farmers who received a payment that was less than the maximum, their effective marginal price is the price offered to them, depending on their sub-treatment group (50 or 100 INR per hour).

To boost precision while avoiding overfitting and weak instruments concerns, we use the instrumental variables LASSO method of Belloni et al. (2012). Our set of candidate instruments consists of indicators for each of the two conservation credit price sub-treatments and their interactions with baseline characteristics. We again choose covariates using double-LASSO, and cluster standard errors by randomization pair.

Intuitively, IV estimates take our ITT estimates and scale them by the fraction of the sample who was in position to respond to the price incentive. Benchmarks were set too low for many farmers to reach, and too high for other farmers, such that they would have reached the maximum payment even without behavior change. The IV estimates instead attribute the full program response to the farmers for whom benchmarks were set appropriately enough to affect their behavior. This method is in the spirit of quasi-experimental estimates of the elasticity of taxable income from non-linear budget sets (as summarized by Saez et al., 2012) and of electricity demand (Ito, 2014).

5.1 Results

Table 4 reports results. Column (1) reports the first-stage relationship for an IV specification with only one instrument: overall eligibility for the conservation credits program. The estimate says that the average effective marginal price in the treatment group was 42 INR per hour.[16] This first-stage relationship is strong, with a very large F-statistic.

Column (2) shows the IV estimate with this one instrument and no covariates, while specifications in columns (3) and (4) add instruments and covariates. Moving across the table, first-stage F-statistics remain strong, while the IV estimates gain precision. Our

---

[16]This value represents a weighted average of the proportion of each group that was marginal, multiplied by the price offered. We separately calculate that 58 percent of farmer-months in the sample faced a positive marginal price.

preferred estimate is in column (4), in which both instruments and covariates are selected by double LASSO. The coefficient of −0.12 implies that average monthly irrigation hours fall by 1 hour for every 8 INR increase in the hourly price. At the middle price of 50, and the control mean of irrigation hours, this implies a price elasticity of 0.13.

One limitation of these IV estimates is that they may overstate the true price elasticity. The exclusion restriction is that the program affected irrigation only through the effective marginal price at the end of the meter reading period. This assumption will be violated if the program affected irrigation for farmers who do not end up facing positive marginal prices in a given month – for example, if they attempted to conserve below the benchmark but failed to reach their target. This is a fundamental limitation of this method for estimating demand.

However, we can still bound the price elasticity using the IV and reduced-form estimates. The IV estimate loads the entire reduced-form effect of the program onto the fraction of farmers with a positive effective marginal price. If some farmers change their behavior but are not observed in this group, then the true proportion of farmers affected by the program is greater than indicated by the first stage. On the other hand, it is unlikely that all farmers in the treatment group were affected by the program, so the true proportion is less than 1. The true price elasticity is then bounded above by the IV estimate, and bounded below by the reduced-form estimate: $(0.16, 0.20)$.[17]

6 Cost-effectiveness

Now, we consider the cost-effectiveness of the conservation credits intervention as implemented in the study, from the perspective of an electricity utility. For now we set aside the social costs of groundwater depletion and focus solely on energy. Suppose political constraints rule out straightforward volumetric prices for electricity. Might a utility company find it less costly to reduce demand via a conservation payments program than to increase supply by procuring additional electricity?

We calculate the cost of reducing electricity demand through this program as the ratio of total expenditures on conservation payments to total energy conserved. Note this ratio is not just a rescaling of our demand estimates, because it includes payments made to inframarginal farmers. For total energy conserved as a result of the program, we use the preferred OLS estimate from Table 2 because it is more precise than the Poisson estimate.

---

[17]Scaling the ITT effect of −11 hours by the average price offered in the treatment group (75 INR per hour) gives a reduced-form effect of −0.15 hr/INR. At the middle price of 50, and the control mean of irrigation hours, this implies a price elasticity of 0.16.

Table 5 shows further details of this calculation, parameters used, and results.

We estimate that the conservation credits intervention reduced electricity use at a cost of 6.1 INR per kWh conserved. This value appears to be similar to published estimates of the costs of electricity procurement. It is slightly greater than the average cost of electricity procurement per unit sold by the electric utility in our study area, 5.4 INR per kWh, but the marginal costs of electricity procurement are likely greater than than average costs. It is also lower than the cost of electricity procurement in the nearby state of Punjab.

In this calculation, we only consider expenditures on conservation payments and omit other program costs such as meter hardware and personnel and travel expenses for reading meters. We do so for two reasons. First, electric utilities obtain other benefits from metering their customers, so we prefer to consider the perspective of a utility that is already collecting this data. Second, metering costs in a permanent program would likely be lower than in our short-term intervention. It would likely be more cost-effective to install smart meters that can be read remotely, saving the labor and travel expenses, the fixed cost of which would then be amortized over a longer period.

It is also worth considering the social costs of groundwater depletion and of air pollution from electricity generation. A utility company may not include these costs in a cost-effectiveness calculation, but a government may want to consider them as motivation for subsidizing a conservation payments program. Estimating the negative externalities from groundwater extraction is beyond the scope of this study. But even relatively small estimates of these social costs would likely make it socially optimal for a utility to offer conservation payments before expanding electricity supply.

7   Optimal contract design

Finally, we consider the optimal choice of benchmark values against which conservation is rewarded. Many conservation programs simply set the benchmark equal to an observed pre-program value of the target outcome – for example, the number of trees on a property at program enrollment, or the amount of water used last year. In some settings, this pre-program value might reliably predict the target outcome absent the program, such that conservation measured relative to this value indeed represents conservation relative to the counterfactual. But in other settings, it might not predict the counterfactual well. And if not, then basing the benchmark on the pre-program value might not be the best approach at all.

We argue that conservation programs like ours should not try to set the benchmark *accurately*; they should instead set the benchmark *cost-effectively*. When the counterfactual

cannot be predicted with precision, one might be tempted to still set the benchmark to the best prediction of the counterfactual. But this is the wrong goal, because it uses the wrong loss function. The cost of prediction error in a conservation program is asymmetric and nonlinear, so optimal benchmarks will not generally coincide with conditional-mean predictions of the counterfactual. We also argue that the cost-effective benchmark value is likely to be stringent. That is, the optimal conservation contract will often reward conservation relative to a benchmark value set substantially lower than the observed pre-program outcome.

The optimal choice of benchmark is a tradeoff between three effects. On one hand, (1) higher benchmarks require more inframarginal payments – the program must reward more units of resource "conservation" that never would have been consumed anyway. On the other hand, (2) higher benchmarks can induce more people to respond to the program. If benchmarks are too low, participants must reduce consumption without compensation before payments kick in, so a low benchmark may not be worth pursuing. However, it is also possible that (3) low benchmarks can induce or motivate participants into conserving more than they otherwise would, through either income effects or behavioral factors.

In our setting, we show that inframarginal payments arising from prediction error are a large share of total payments, implying they are a first-order concern for program cost-effectiveness. We were able to reduce program costs by setting benchmarks lower than predicted counterfactuals, especially for participants with high expected variance. We also show that many participants do respond to aspirationally low benchmarks by conserving more – yielding a seemingly "free lunch" of both lower costs and greater benefits. These empirical results suggest that the optimal benchmarks are generally lower than either the pre-program value or the counterfactual. Finally, we estimate the overall cost-effectiveness of the program in the high- and low- benchmark subtreatment groups. On average, conservation in the low-benchmark group was only 41 percent as expensive as in the high-benchmark group.

7.1   Low benchmarks reduce inframarginal payments

First, we show that inframarginal payments arising from prediction error are large in our setting. Inframarginal payments – payments for not using units of a resource that would not have been used without the program – mechanically rise with the benchmarks. The higher the benchmark, the more units of conservation are measured and rewarded. Therefore, program costs can be reduced by setting lower benchmarks, ignoring any behavioral response for now.

19

Inframarginal payments can be estimated by calculating what payments would have been in the control group under alternative benchmark schemes. There is no behavioral response in the control group, so if benchmarks were set equal to perfect predictions of pumping, these calculated payments would be zero.

Table 13 lists average inframarginal payments calculated under each of several alternative benchmark scenarios. Inframarginal payments under all benchmark scenarios are high relative to average realized payments in the treatment group. In fact, these calculations imply that most of the payments disbursed in our program were inframarginal. Inframarginal payments also vary widely based on how benchmarks are set. For example, setting benchmarks equal to the pre-program value would have nearly doubled payment expenditures.

Who should get low benchmarks relative to their expected counterfactual – everyone, or only some people? We suggest that cost-effectiveness can likely be improved by targeting benchmarks on expected variance of the counterfactual (or, the expected magnitude of the prediction error). The case for lowering benchmarks is stronger when the counterfactual is highly uncertain, and weak when the counterfactual can be predicted with precision.

In our program, pumping variance increases with observed pre-program pumping, and thus so does prediction error and inframarginal payments.[18] For participants whose pre-program pumping exceeds 200 or 300 hours, setting benchmarks close to their predicted counterfactual may not be worthwhile. Some of these participants end up pumping near zero during the program even absent incentives, so they would receive expensive inframarginal payments. Depending on the behavioral response, it may be cost-effective to set low benchmarks for people with high pre-program pumping, essentially giving up on most of them.

This logic inspired the design of the benchmarks that we implemented in the treatment group. Figure 7 plots benchmarks as a function of pre-program pumping, with the $y = x$ line for comparison. Benchmarks start low at low values of pre-program pumping, rise roughly one-for-one for a range of values, but then flatten out and decline again, such that many participants with very high pre-program pumping faced the minimum benchmark of 10 hours. The broad shape of this function came from optimizing a nonlinear function of pre-program pumping to maximize the share of participants who would be marginal to the incentive (i.e., pump at or below the benchmark) subject to the project's budget constraint. This function was increased or decreased by 15 percent in the high-benchmark

---

[18]Appendix Figure 9 clearly shows this pattern of heteroskedasticity, and column 4 of Appendix Table 12 confirms it by regressing log-squared residuals on pre-program pumping.

and low-benchmark randomized sub-treatment groups. Benchmarks were rounded to the nearest 10 hours, set separately for each well, and summed. Remaining variation comes from a small number of errors in field operations, as benchmarks were calculated during the same visit at which pre-program pumping was observed.

7.2 Low benchmarks can induce more conservation

Now, we turn to the behavioral response. How does pumping respond to a higher or lower benchmark? Lower benchmarks reduce inframarginal payments, reducing program costs. But if benchmarks are too low, participants might choose not to conserve at all, reducing program benefits as well.

However, there is another possible effect. Benchmarks set low, but not too low, might be able to induce participants to conserve more than they would with higher benchmarks. This effect could operate through income effects: when participants must cut back more in order to get paid, they very well might choose to do so. It could also operate through behavioral channels, if participants put more weight on the cash rewards than on the costs of forgone resource consumption, or if they interpret the benchmark as containing information about scientific recommendations or social norms. (This could occur even without the program making any such claims; ours did not.)

To find out which effect dominates, we can study our randomized benchmark sub-treatment. Benchmarks were nudged 15 percent higher (before rounding) for half of participants and 15 percent lower for the other half. The first stage of this sub-treatment is strong: Table 6, column 1 shows that it increased benchmarks by 9 hours on average with an F-statistic of 135.

Overall evidence is not strong that random marginal changes to benchmarks affect pumping. In Table 6, column 2 reports a regression of pumping hours on overall treatment (i.e., payment eligibility) and its interaction with the high-benchmark sub-treatment. Relative to the low-benchmark group, the high-benchmark group pumped about 5 hours more per month, but the confidence interval extends from −1 to 10.

However, there are strong reasons to expect the effect of benchmarks to differ across participants. At low values of water use, lowering benchmarks is unlikely to increase conservation. Both benchmarks and water use are bounded by zero, so both the available conservation and the available rewards are limited. At high values of water use, it is much more plausible that we might see an "aspirational benchmarks" phenomenon.

This is in fact what we find: Participants with high expected pumping respond to lower benchmarks by conserving more. Columns 3-4 of Table 6 report the same regres-

sion estimated in "low pumping" and "high pumping" subsamples, which split the full sample at 60 hours of pre-program pumping. (This threshold is approximately where our benchmark formula stops rising linearly and flattens out.) For participants with low pre-program pumping, the benchmark has no effect. But for participants with high pre-program pumping, there is a sizable positive effect – lower benchmarks reduce pumping. Columns 5-6 report 2SLS versions of the regressions; at least within the range of benchmarks we set, lower benchmarks appear to reduce pumping roughly one-for-one.

This pattern holds when we expand the analysis. Figure 8 plots the effect of the high-benchmark sub-treatment by seven quantiles of pre-program pumping. These estimates come from a Poisson regression so that the effects are comparable to each other in proportional terms. For participants with low pre-program pumping, a higher benchmark decreases pumping by around 20 percent. For participants with high pre-program pumping, a higher benchmark increases pumping by 20 to 30 percent. These results imply that to maximize water conservation, the program should offer higher benchmarks for participants with low expected pumping (so that conservation is worthwhile for more of them) and lower benchmarks for participants with high expected pumping (because they can be induced to conserve more that way).

7.3 Low benchmarks can improve cost-effectiveness

Finally, we put the inframarginal payments and behavioral response together and calculate the overall cost-effectiveness of the program under different scenarios. We do not attempt to globally optimize the benchmarks, which would require considerable out-of-sample extrapolation. Instead, we restrict attention to the local variation introduced by our randomized high- and low-benchmark sub-treatment groups.

We simulate locally-optimized benchmarks by choosing either the high or low benchmarks to maximize conservation, using the quantile estimates from Figure 8. Within each quantile bin, we keep participants in the high-benchmark group if the effect of higher benchmarks on pumping is negative, keep participants in the low-benchmark group otherwise, and drop participants in the other group. Using this carefully selected sub-sample, we then recalculate overall program effects (i.e., relative to the full control group) and payment expenditures.

Results are shown in Table 7. Expenditures on payments are similar in the conservation-maximizing group as in the full sample, even though it achieves more conservation. As a result, the average cost of conservation is only 86 INR per hour in the conservation-maximizing group, as compared with 117 INR per hour in the full sample. Simply locally

22

optimizing benchmarks within the range of our experiment can produce a 27 percent increase in cost-effectiveness.

In comparison, the high-benchmark group is much less cost-effective. It achieves less conservation while disbursing more payments, making it 47 percent more expensive than the full sample and almost exactly twice as expensive as the conservation-maximizing group.

This simulation only maximizes conservation, not cost-effectiveness. As it turns out, the low-benchmark group is even more cost-effective than the conservation-maximizing group. It reduces pumping at a cost of 71 INR per hour, which is 39 percent less expensive than the full sample. Jointly optimizing both conservation and payment expenditures would likely produce a schedule of benchmarks that is even more cost-effective.

8   Conclusion

This study finds that moderately sized incentives for groundwater conservation lead farmers to reduce groundwater irrigation by approximately 20 percent. Impacts increase over time, indicating that the response to incentives can be sustained. And this is a short-term response: Our program lasted for only one irrigation season and was introduced after crops were already planted. In a longer-term program, the response would likely be even greater, since farmers would be able to substitute crops and adjust other inputs.

These findings suggest that conservation credits, a policy solution similar to existing "payments for environmental services" programs, are an effective tool for managing groundwater and energy resources in India. In many settings — and perhaps especially in the setting of agricultural groundwater extraction in India — Pigouvian taxes may be politically infeasible. By exchanging corrective taxes for subsidies, conservation credits overcome the political barriers to taxing the agricultural sector, while still introducing marginal incentives for conservation. Thus, conservation credits may be a particularly promising policy approach for reducing inefficient groundwater extraction.

We also find that the program effect is large relative to total cost of incentives: as designed, the overall expenditure per unit of energy conserved is similar to the per-unit cost a utility company would face in procuring electricity. This suggests that a utility capable of rolling out conservation credits at low fixed cost could potentially save money if the program were carefully designed. Our program uses a combination of individual-specific benchmarks (set using verifiable baseline irrigation information) and maximum payments to avoid extreme payments for infra-marginal behavior. Yet our program design leaves room for further improvements in benchmark targeting and careful setting of

maximum payments. Understanding how to optimally set benchmarks and maximum payments, including understanding the best information to use for setting these parameters, is a key question for future work.

Figure 1: Villages in Study Area

Notes: This figure shows the villages in Gujarat, India where participants were enrolled as blue dots.
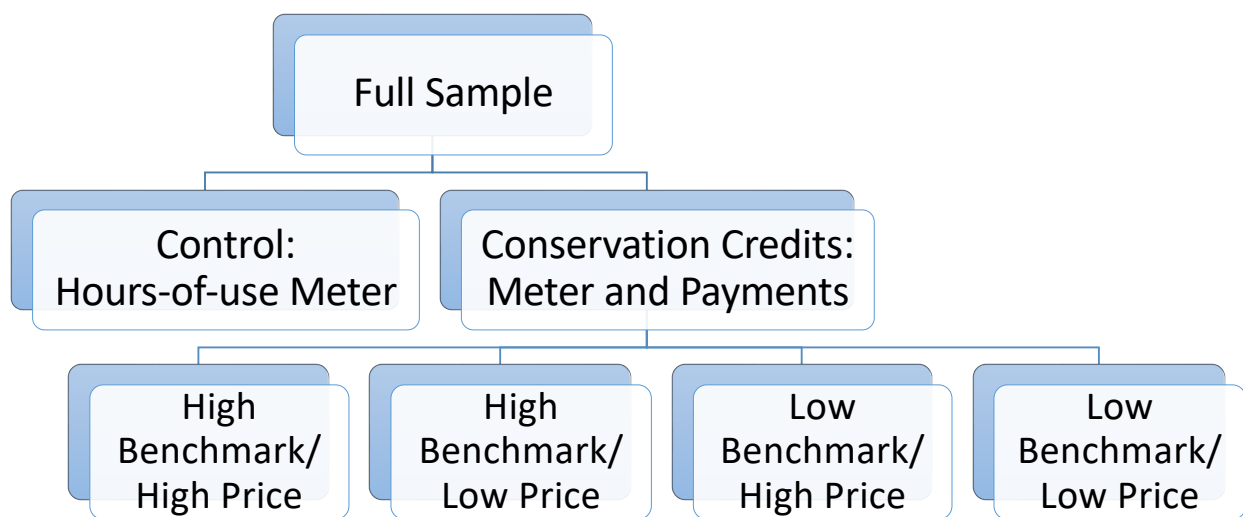
Figure 2: Intervention Design

Notes: This figure illustrates the four interventions used in the randomized experiment. Farmer sharing pools were assigned in equal proportion to the Control and Conservation Credits groups. Within the Conservation Credits group, the four sub-treatments were assigned in equal proportion.

Sample Frame	Partner Lists
Summer 2022	Eligibility, Baseline and Consent
Fall 2022	Meter Installation
Dec 2022	Baseline Meter Reading, Randomization
Jan 2023	1st Meter Reading Follow-up
Feb 2023	2nd Meter Reading Follow-up
Mar 2023	3rd Meter Reading Follow-up
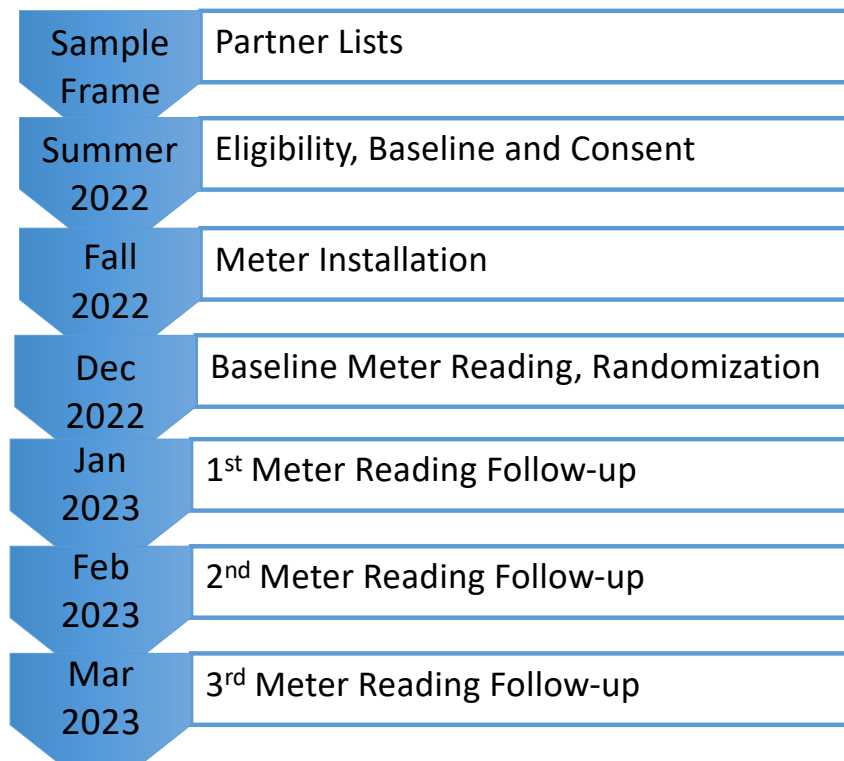
Figure 3: Experiment Timeline

Notes: This figure displays the timeline of our experimental intervention and data collection processes.

(a) Hours of irrigation

(b) Hours relative to benchmark

Figure 4: Bin treatment effects of payment eligibility on the distribution of irrigation hours

Notes: This figure plots estimated coefficients from regressions of binary indicators for whether a participant's value of hours of irrigation fell within specified ranges on eligibility for the Conservation Credits program. Regressions use a subset of controls. Error bars represent 95% confidence intervals.

(a) OLS *(b) Poisson*

Figure 5: Treatment effects of conservation credits grew over time

Notes: This figure plots the treatment effect of offering Conservation Credits on hours of irrigation across the three months of the intervention period. Treatment effects are estimated using double-LASSO selected controls. Error bars represent 95% confidence intervals.

Figure 6: Budget set of conservation credits.

This figure shows the general form of the budget set created by a conservation credit program, along with indifference curves of two representative participants. The payment equals the price p times the quantity units conserved below the benchmark, up to a maximum payment. Irrigator A is marginal and will respond to the program by reducing quantity extracted. Irrigator B is extra-marginal, and does not change quantity extraction in response to the program.

29

Figure 7: Benchmarks in the Treatment Group

Notes: Figure plots benchmarks set in the treatment group. Points are jittered to show density; all benchmarks were multiples of 10 hours.

Figure 8: Effects of High Benchmarks on Pumping, By Quantiles of Pre-Program Pumping

Notes: Estimated heterogeneous effects of the randomized high-benchmark sub-treatment group relative to the low-benchmark group, by quantiles of hours of irrigation before the intervention. Figure plots coefficients in Poisson regression.

Table 1: Baseline Summary Statistics in Full Sample and by Treatment Group

|	Full Sample		Control	Treatment
	(1)	(2)	(3)	(4)
	Mean	SD	Mean	Mean
A. Demographics				
Household size	6.34	2.85	6.46	6.22
Scheduled caste/tribe or other backwards caste	0.86	0.34	0.86	0.87
Muslim	0.09	0.28	0.09	0.09
Years of education (household head)	10.94	3.39	10.88	11.00
Literacy (household head)	0.82	0.38	0.83	0.81
B. Farm statistics				
Plot hectares	1.95	1.35	1.97	1.92
Number of crops cultivated	1.96	1.08	2.01	1.91
Fraction of farmed area planted with cotton	0.53	0.41	0.54	0.53
Fraction of farmed area planted with sorghum/millet	0.15	0.25	0.15	0.16
Fraction of farmed area planted with groundnut	0.15	0.25	0.14	0.15
Fraction of farmed area planted with pulses	0.11	0.21	0.11	0.10
Has cow, ox, or buffalo	0.92	0.27	0.93	0.91
Has plow or tractor	0.50	0.50	0.50	0.50
C. Well Statistics				
Total number of active wells	1.19	0.39	1.19	1.19
Deepest well is dugwell	0.24	0.43	0.23	0.26
Deepest well is borewell	0.25	0.43	0.23	0.27
Deepest well is dug-cum-borewell	0.51	0.50	0.55	0.47
Deepest well: ever deepened	0.17	0.37	0.17	0.17
Deepest well: depth (meters)	58.62	85.17	53.66	63.12
Deepest well: max water level (meters)	16.07	36.60	14.68	17.33
Deepest well: pump power	5.61	3.27	5.46	5.75
D. Irrigation Statistics				
Pre-intervention monthly irrigation hours	71.71	71.09	69.81	73.43
Total self-reported hours of irrigation on farm	340.97	2205.91	327.45	353.25
Total self-reported hours of irrigation off farm	32.46	153.97	32.17	32.73
Purchased water for irrigation	0.01	0.11	0.01	0.01
Used drip irrigation	0.41	0.49	0.42	0.41
Used sprinkler irrigation	0.01	0.10	0.01	0.02
Used raised beds	0.69	0.46	0.69	0.68
Used rotational, strip, or zero-tillage	0.19	0.39	0.17	0.20
Used farm bunds	0.09	0.29	0.10	0.08
Test for joint orthogonality of covariates				
F-statistic			0.64	
P-value			0.93	
Sample size				
Number of individuals	989		471	518
Percent of sample	100.0		47.6	52.4

Notes: This table summarizes baseline characteristics of the sample of farmers who completed all three meter reading survey rounds during the intervention. The F-statistic and associated P-value test the joint orthogonality of all characteristics listed in the table to treatment assignment relative to the control group.

Table 2: *Intent-to-Treat Impacts of Conservation Credits on Hours of Irrigation*

	OLS				Poisson			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Payment Eligibility	-5.97*	-9.92***	-10.4***	-9.01***	-0.14*	-0.23***	-0.25***	-0.22***
	[3.46]	[3.21]	[2.44]	[2.92]	[0.079]	[0.073]	[0.056]	[0.067]
Payment Eligibility × High Price				-2.80				-0.047
				[3.48]				[0.084]
Control Mean	46.59	46.59	46.59	46.59	46.59	46.25	46.59	46.59
Village FE		X				X		
Lasso Controls			X	X			X	X
N Clusters	494	494	494	494	494	485	494	494
N Farmers	989	989	989	989	989	970	989	989
N Observations	2,967	2,967	2,967	2,967	2,967	2,910	2,967	2,967

Notes: The sample includes all farmers who completed all three meter reading survey rounds during the intervention. The outcome is monthly hours of irrigation by the farmer during the three intervention months (scaled to 31 days). Standard errors clustered at the randomization pair level are in brackets.

Table 3: *Intent-to-Treat Impacts of Conservation Credits on Energy Use (kWh)*

	OLS				Poisson			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Payment Eligibility	-23.5	-98.4	-139.8***	-83.7	-0.039	-0.17	-0.22***	-0.14*
	[70.2]	[64.5]	[40.6]	[53.9]	[0.12]	[0.11]	[0.065]	[0.082]
Payment Eligibility × High Price				-111.2*				-0.16
				[66.4]				[0.11]
Control Mean	610.81	610.81	610.81	610.81	610.81	606.16	610.81	610.81
Village FE		X						
Lasso Controls			X	X			X	X
N Clusters	494	494	494	494	494	485	494	494
N Farmers	989	989	989	989	989	970	989	989
N Observations	2,967	2,967	2,967	2,967	2,967	2,910	2,967	2,967

Notes: The sample includes all farmers who completed all three meter reading survey rounds during the intervention. The outcome is monthly kWh of energy used for irrigation by the farmer during the three intervention period months (scaled to 31 days). Energy use is calculated from hours of irrigation as described in Section 2. Standard errors clustered at the randomization pair level are in brackets.

Table 4: Demand for Groundwater Irrigation: Instrumental Variables Estimation

	First Stage	IV		
	(1)	(2)	(3)	(4)
Marginal Price		-0.14*	-0.13*	-0.12**
(INR/Hour)		[0.080]	[0.078]	[0.049]
Payment	42.2***			
Eligibility	[1.44]			
Outcome Control Mean	0.00	46.59	46.59	46.59
CD Wald F-stat		1,562.17	900.71	17.35
Fixed Effects		Month	Month	LASSO
Controls				LASSO
Instruments		Treatment	Price Sub-Treatments	LASSO
N Instruments		1	2	7
N Clusters		494	494	494
N Farmers	989	989	989	989
N Observations	2,967	2,967	2,967	2,967

Notes: The sample includes farmer-months among farmers who remained in the experiment until the final meter reading. The outcome is the monthly hours of irrigation in each of the three intervention period survey rounds. The marginal price of an hour of irrigation is instrumented using the the overall Conservation Credit treatment in Columns 2, using all four Conservation Credit sub-treatments in Column 3, and using additional high-dimensional instruments selected by double-LASSO in Column 4. Standard errors clustered at the randomization pair level are in brackets.

Table 5: Cost-Effectiveness of Conservation Payments

Parameter	Value	Unit	Source
Panel A: Parameters used			
Pump motor efficiency, from a similar context	40%	-	Mitra, Balasubramanya, & Brouwer (2023)
Unit conversion constant	0.7457	kW per hp	Known constant
Mean duration of intervention	3.7	Months	Meter reading data
Panel B: Calculation of cost-effectiveness			
Average effect of program on electricity use, monthly	-139.8	kWh/month per farmer	Table 3, column (3)
Average effect of program, scaled to rabi season	-512.3	kWh per farmer	Calculated
Average conservation payments, rabi season	3369	INR per farmer	Program implementation data
Average expenditure per unit electricity conserved	6.6	INR/kWh	Calculated
Panel C: Comparisons of cost-effectiveness			
Cost of reducing electricity use through this program	6.6	INR/kWh	From above
Average cost of electricity procurement per unit sold, Gujarat	5.4	INR/kWh	Paschim Gujarat Vij Company Ltd. (2021)
Cost of electricity procurement, Punjab	7.9	INR/kWh	Mitra, Balasubramanya, & Brouwer (2023)

Table 6: Effects of Benchmarks on Hours of Irrigation, by Pre-Program Pumping

	First Stage	Reduced Form			IV	
	(1) All	(2) All	(3) Low	(4) High	(5) Low	(6) High
Payment Eligibility		-10.1*** [2.79]	-2.64 [2.61]	-20.6*** [6.09]		
Eligibility x High Benchmark	9.09*** [0.78]	4.58 [2.97]	-2.80 [2.42]	15.1** [6.69]		
Benchmark (hours)					-0.64 [0.43]	1.10** [0.47]
Control Mean	0.00	46.59	18.69	85.39	0.00	0.00
Pre-program pumping (fine bin FE)	X	X	X	X	X	X
N Clusters	265	492	284	208	154	111
N Farmers	518	989	573	416	299	219
N Observations	1,554	2,967	1,719	1,248	897	657

Notes: "Low" and "High" samples split the full sample at 60 hours of pre-program pumping. The full sample includes all farmers who completed all three meter reading survey rounds during the intervention. The outcome is monthly hours of irrigation by the farmer during the three intervention months (scaled to 31 days). Standard errors clustered at the randomization pair level are in brackets.

Table 7: Program Cost-Effectiveness by Benchmark Group

Metric	Full sample	High-benchmark group	Low-benchmark group	Locally conservation-maximizing benchmarks
Average effect of program on pumping, rabi season (hours)	-28.7	-36.6	-24.0	-40.3
Average conservation payments, rabi season (INR)	3,369	2,605	4,157	3,482
Average expenditure per unit of pumping conserved (INR/hour)	-118	-71	-173	-86

References

Ahrens, A., Hansen, C., and Schaffer, M. E. (2018). pdslasso and ivlasso: Progams for post-selection and post-regularization ols or iv estimation and inference. *Statistical Software Components, Boston College Department of Economics*.

Asher, S. and Novosad, P. (2020). Rural roads and local economic development. *American Economic Review*, 110:797–823.

Athey, S. and Imbens, G. W. (2017). The state of applied econometrics: Causality and policy evaluation. volume 31, pages 3–32. American Economic Association.

Bai, Y., Hsieh, M. H., Liu, J., and Tabord-Meehan, M. (2023). Revisiting the analysis of matched-pair and stratified experiments in the presence of attrition. *Journal of Applied Econometrics*.

Belloni, A., Chen, D., Chernozhukov, V., and Hansen, C. (2012). Sparse Models and Methods for Optimal Instruments with an Application to Eminent Domain. *Econometrica*, 80(6):2369–2429.

Belloni, A., Chernozhukov, V., and Hansen, C. (2013). Inference on treatment effects after selection among high-dimensional controls. *Review of Economic Studies*, 81(2):608–650.

Blakeslee, D., Fishman, R., and Srinivasan, V. (2020). Way down in the hole: Adaptation to long-term water loss in rural india. *American Economic Review*, 110(1):200–224.

Börner, J., Baylis, K., Corbera, E., Ezzine-de Blas, D., Honey-Rosés, J., Persson, U. M., and Wunder, S. (2017). The Effectiveness of Payments for Environmental Services. *World Development*, 96:359–374.

Bruno, E. M. and Jessoe, K. (2021). Missing markets: Evidence on agricultural groundwater demand from volumetric pricing. *Journal of Public Economics*, 196.

Bruno, E. M., Jessoe, K. K., and Hanemann, W. M. (2024). The dynamic impacts of pricing groundwater. *Journal of the Association of Environmental and Resource Economists*, 11(5):1201–1227.

Burgess, R., Greenstone, M., Ryan, N., and Sudarshan, A. (2020). The consequences of treating electricity as a right. *Journal of Economic Perspectives*, 34(1):145–69.

Chakravorty, U., Dar, M. H., and Emerick, K. (2023). Inefficient water pricing and incentives for conservation. *American Economic Journal: Applied Economics*, 15(1):319–50.

Chen, J. and Roth, J. (2024). Logs with zeros? some problems and solutions. *Working Paper*.

de Chaisemartin, C. and Ramirez-Cuellar, J. (2024). At what level should one cluster standard errors in paired and small-strata experiments? *American Economic Journal: Applied Economics*, 16:193–212.

Fishman, R., Lall, U., Modi, V., and Parekh, N. (2016). Can Electricity Pricing Save India's Groundwater? Field Evidence from a Novel Policy Mechanism in Gujarat. *Journal of the Association of Environmental and Resource Economists*, 3(4):819–855.

Gonzalez-Alvarez, Y., Keeler, A. G., and Mullen, J. D. (2006). Farm-level irrigation and the marginal cost of water use: Evidence from Georgia. *Journal of Environmental Management*, 80(4):311–317.

Hendricks, N. P. and Peterson, J. M. (2012). Fixed Effects Estimation of the Intensive and Extensive Margins of Irrigation Water Demand.

Ito, K. (2014). Do Consumers Respond to Marginal or Average Price? Evidence from Nonlinear Electricity Pricing. *American Economic Review*, 104(2):537–563.

Jack, B. K. and Cardona Santos, E. (2017). The leakage and livelihood impacts of PES contracts: A targeting experiment in Malawi. *Land Use Policy*, 63:645–658.

Jack, B. K., Jayachandran, S., Kala, N., and Pande, R. (2023). Money (not) to burn: Payments for ecosystem services to reduce crop residue burning. *Working Paper*.

Jack, B. K. and Smith, G. (2016). Charging Ahead: Prepaid Electricity Metering in South Africa. *NBER Working Paper Series*.

Jasechko, S., Seybold, H., Perrone, D., Fan, Y., Shamsudduha, M., Taylor, R. G., Fallatah, O., and Kirchner, J. W. (2024). Rapid groundwater decline and some cases of recovery in aquifers globally. *Nature*, 625:715–721.

Jayachandran, S., De Laat, J., Lambin, E. F., Stanton, C. Y., Audy, R., and Thomas, N. E. (2017). Cash for carbon: A randomized trial of payments for ecosystem services to reduce deforestation. *Science*, 357(6348):267–273.

Lee, D. S. (2009). Training, wages, and sample selection: Estimating sharp bounds on treatment effects. *Review of Economic Studies*, 76(3):1071–1102.

Lee, K., Miguel, E., and Wolfram, C. (2018). Experimental Evidence on the Economics of Rural Electrification. *Working Paper*.

McRae, S. (2015). Infrastructure Quality and the Subsidy Trap. *American Economic Review*, 105(1):35–66.

Meenakshi, J., Banerji, A., Mukherji, A., and Gupta, A. (2013). Does marginal cost pricing of electricity affect groundwater pumping behaviour of farmers? Evidence from India. *3ie Impact Evaluation Report*, 4.

Mitra, A., Balasubramanya, S., and Brouwer, R. (2023). Can cash incentives modify groundwater pumping behaviors? evidence from an experiment in punjab. *American Journal of Agricultural Economics*, 105:861–887.

Mullahy, J. and Norton, E. C. (2023). Why transform y? the pitfalls of transformed regressions with a mass at zero*. *Oxford Bulletin of Economics and Statistics*.

Pattanayak, S. K., Wunder, S., and Ferraro, P. J. (2010). Show me the money: Do payments supply environmental services in developing countries? *Review of Environmental Economics and Policy*, 4(2):254–274.

Pfeiffer, L. and Lin, C.-Y. C. (2014). The effects of energy prices on agricultural groundwater extraction from the high plains aquifer. *American Journal of Agricultural Economics*, 96(5):1349–1362.

Registrar General and Census Commissioner of India (2001). Census of India 2001.

Ryan, N. and Sudarshan, A. (2022). Rationing the commons. *Journal of Political Economy*, 130:210–257.

Saez, E., Slemrod, J., and Giertz, S. H. (2012). The Elasticity of Taxable Income with Respect to Marginal Tax Rates: A Critical Review. *Journal of Economic Literature*, 50(1):3–50.

Sekhri, S. (2014). Wells, water, and welfare: The impact of access to groundwater on rural poverty and conflict. *American Economic Journal: Applied Economics*, 6(3):76–102.

Silva, J. M. C. S. and Tenreyro, S. (2006). The log of gravity.

Smith, S. M., Andersson, K., Cody, K. C., Cox, M., and Ficklin, D. (2017). Responding to a groundwater crisis: The effects of self-imposed economic incentives. *Journal of the Association of Environmental and Resource Economists*, 4(4):985–1023.

Sovacool, B. K. (2017). Reviewing, Reforming, and Rethinking Global Energy Subsidies: Towards a Political Economy Research Agenda. *Ecological Economics*, 135:150–163.

A Covariates included in LASSO-Selection

The following variables, which are collected prior to randomization, are fed into each double-selection LASSO after being interacted with 2 survey visit (i.e., month) indicators:

- Meter Reading

    - Average daily hours of pumping in the first month

    - Average daily energy use in the first month (calculated)

    - The natural logs of the above two variables (imputed to zero if the argument is 0)

- Baseline Survey

    - 121 Village indicators

    - Total wells on the primary farm

    - Whether deepest well has ever been deepened

    - Water level on deepest well

    - Whether deepest well went dry the previous Kharif season

    - Depth of deepest well

    - Indicators for deepest well being borewell or dug-cum-borewell (dugwell is omitted well type)

    - Pump power for pump on deepest well

    - Number of crops cultivated

    - Indicator for above-median number of crops cultivated

    - Fraction of farmed area planted with each of: cotton, sorghum or millet, groundnut, and pulses

    - Total self-reported hours of irrigation (a) on primary farm and (b) off primary farm, previous Kharif season

    - Indicator for whether purchased water for irrigation during previous Kharif season

    - Indicators for typical use of raised beds, farm bunds, and low/zero-tillage practices (common water conservation practices) and for use drip irrigation previous Kharif

- Years of education (household head)

- Above-median years of education (household head)

- Indicators for Hindu and Muslim (omitted religion is other)

- Indicator for Scheduled Caste/Scheduled Tribe/Other Backwards Caste

- Groundwater Prospects Maps

  - The distance to the nearest recharge structure, fracture, escarpment, water body, observation well, mapped dugwell, mapped handwell, mapped borewell, canal, dyke, railway, stream, and road and their natural logs

  - The number of recharge structures, fractures, canals, and streams and the length of fractures in 1, 2, 3, 4 and 5 km radii

  - The fraction of land area made up of each of 42 rock types in 1, 2, 3, 4 and 5 km radii

- Satellite Images

  - The natural log of the seasonal differences in EVI and NDVI in the 2021-22 rabi season

  - the mean EVI and NDVI in November 2022

B   Robustness to Attrition

Appendix Table 8 reports the response rates by treatment status and meter reading visit. While randomization was assigned at baseline, neither the farmer nor the surveyor was aware of treatment assignment until the end of the first visit shown (i.e., the "Randomization visit"). Thus, attrition at the first visit is not due to random assignment. However, following randomization, attrition increases at each subsequent meter reading visit, and differentially more so for the control group.

We examine the robustness of our results to attrition in two ways. First, we examine evidence that different types of people left the experiment in treatment and control. Second, we show that the Lee bounds on the impact of the conservation credits program on hours of irrigation do not substantially reduce our estimates.

We begin by examining the evidence on whether attrition led to baseline imbalance across treatment. Table 1 in the text shows that the treatment and control farmers are well balanced on observable baseline characteristics (p-value for joint orthogonality test

is 0.93). Second, we focus on the characteristics of those who attrit after randomization. Appendix Table 9 shows that we cannot reject that these attriters' baseline characteristics are jointly orthogonal to treatment status (we also focus specifically on differential attrition by baseline irrigation hours in this table, and find no evidence on this dimension). These checks show that there are not observable differences between the treatment and control groups either among those who attrit or among those who remain in the the sample.

We next turn to the possibility that unobservable differences in attritors are driving our treatment effects. Specifically, we follow Lee (2009) in bounding the treatment effect of conservation credits on hours of irrigation under a monotonicity assumption: that no one who attrited in control would have been *more* likely to attrit in treatment. Because our outcome data are a panel, we remove the respondents in the treatment group with the highest (lowest) average hours of irrigation across the intervention period and then re-estimate Equation 3 without controls to find the lower (upper) bound of the treatment effect. We trim the treatment group by 7.2%, which is the difference in attrition rates (6.7%) between the two groups as a proportion of the retention rate between the randomization visit and the fourth meter reading in the treatment group (92%). Appendix Table 10 shows that even the upper bound on the treatment effect is negative and less than three hours different from the ITT estimate, although it not statistically different from zero (bounding also sacrifices statistical precision because rich controls cannot be included). The upper bound is correct only if differential attrition is driven by the control farmers who irrigate the fewest hours. However, Table 9 shows that baseline irrigation hours tend to be lower among retained control farmers but not among treatment farmers, suggesting that control farmers who irrigate more than average hours are driving differential attrition (baseline irrigation is strongly positively correlated with intervention period irrigation). Thus, our ITT estimates are likely to be biased upward in the positive direction. In fact, we find that the lower bound of the conservation credits impact – which is correct if the control farmers who attrit are the highest irrigators – is -19.2 hours, more than 13 hours below the ITT estimate. Overall, these bounds demonstrate that it is very unlikely that differential attrition is driving the result that conservation credits reduce groundwater consumption.

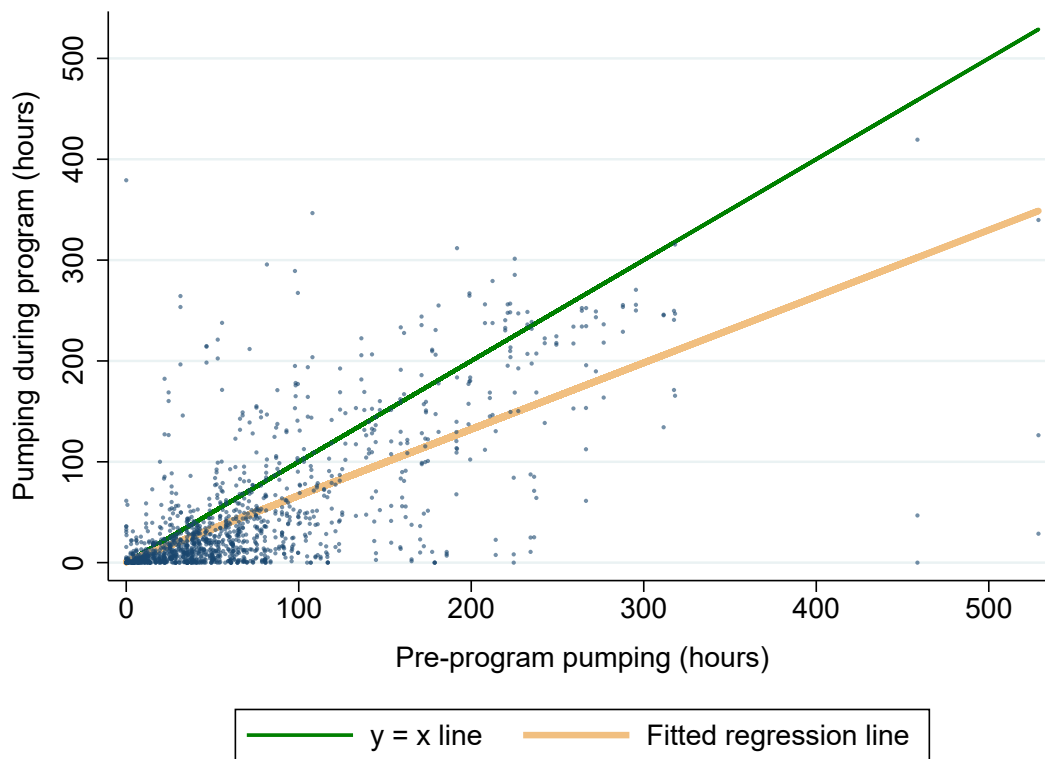# C   Supplementary Tables and Figures

Figure 9: Predicting Pumping in the Control Group

Notes: Monthly hours of irrigation plotted against hours of irrigation before the intervention, for the control group. All values are scaled to a 31-day rate.

Table 8: Sample Retention in Treatment and Control Groups over Time

	Response rates in control group	Difference between treatment and control
	(1)	(2)
A. Response rates		
Randomization visit	0.844	0.012
	(0.364)	(0.020)
		[0.557]
Meter reading 1	0.764	0.052**
	(0.425)	(0.023)
		[0.022]
Meter reading 2	0.738	0.064***
	(0.440)	(0.023)
		[0.006]
Meter reading 3	0.722	0.067***
	(0.448)	(0.024)
		[0.005]
N	673	1308

Notes: This table summarizes attrition in the control and conservation credits treatment after baseline at each of four visits: the initial intervention visit (before randomization), and then each of three meter reading visits. The first column reports the mean and standard deviation for the control group of the farmers reached at each visit shown in the left column. Column (2) reports the estimated differences between treatment and control in the fraction of farmers reached at each visit shown in the left column. The coefficients are from a regression of an indicator from being reached on an indicator of being randomly assigned to the conservation credits treatment with no controls. Standard errors clustered at the randomization pair level are in parentheses; per comparison *P*-values are in square brackets.

Table 9: Baseline covariates do not predict differential retention rates in treatment and control

	(1)	(2)	(3)
Conservation Credits	0.067***	0.039	0.13
	[0.018]	[0.028]	[0.12]
Baseline irrigation hours		-0.00046*	-0.00037
		[0.00024]	[0.00027]
Baseline irrigation hours		0.00038	0.00035
× Conservation Credits		[0.00031]	[0.00036]
Control Mean	0.86	0.86	0.86
F-statistic		1.46	1.26
P-value		0.23	0.19
Baseline Controls x Treatment			X
N Clusters	541	541	540
N Farmers	1,111	1,111	1,110

Notes: This table summarizes how attrition in the control and conservation credits treatment depends on baseline covariates. Each column reports coefficient estimates from regressions of an indicator from being reached at the third meter-reading visit (i.e., being in the analysis sample) on treatment assignment. In columns 2-4, we include covariates and their interaction with treatment assignment: column 1 includes baseline survey covariate, column 2 includes irrigation hours in the month prior to randomization (available for part of the sample), and column 3 includes both baseline covariates and irrigation hours. We report the F-statistics and P-values from a test of the joint significance of all included covariates interacted with treatment assignment: in all specifications, covariates fail to differentially predict retention across treatment assignment. Standard errors clustered at the randomization pair level are in brackets.

Table 10: Lee Bounds on Impact of Conservation Credits Treatment

	ITT Estimate	Upper Bound	Lower Bound
	(1)	(2)	(3)
Payment Eligibility	-5.97*	-2.77	-19.2***
	[3.46]	[3.53]	[3.06]
Control Mean	46.59	46.59	46.59
N Clusters	495	479	471
N Farmers	989	951	951
N Observations	2,967	2,853	2,853

Notes: This table shows Lee bounds on the main treatment effect. Each column reports coefficient estimates from regressions of an indicator from being reached at the third meter-reading visit (i.e., being in the analysis sample) on treatment assignment. Standard errors clustered at the randomization pair level are in brackets.

Table 11: Intent-to-Treat Impacts of Conservation Credits on Yields

	EVI				NDVI			
	(1) Log Diff.	(2) Diff.	(3) Mean	(4) Monthly	(5) Log Diff.	(6) Diff.	(7) Mean	(8) Monthly
Payment Eligibility	0.076 [0.047]	0.20 [0.17]	-0.023 [0.039]	0.0089 [0.026]	0.066 [0.051]	0.0063 [0.0048]	0.00073 [0.0034]	0.0089 [0.026]
Control Mean	-0.58	0.89	1.07	0.96	-2.97	0.07	0.25	0.96
Lasso Controls	X	X	X	X	X	X	X	X
N Clusters	615	615	615	488	615	615	615	488
N Farmers	973	973	973	973	973	973	973	973
N Observations	973	973	973	2,919	973	973	973	2,919

Notes: The sample includes all farmers who completed all three meter reading survey rounds during the intervention. The outcomes include various transformations of EVI and NDVI during the 2022-23 rabi season, including: the natural log of the maximum value less the average of the first four weeks; the maximum value less the average of the first four weeks; the seasonal mean; and the monthly mean. Standard errors clustered at the randomization pair level are in brackets.

Table 12: Predicting Pumping in the Control Group

	Hours of irrigation			Log-squared residuals
	(1)	(2)	(3)	(4)
Pre-program pumping (hours)	0.66*** [0.054]			0.020*** [0.0017]
Baseline covariates (basic set)		X		
Baseline covariates (rich set)			X	
Block FE		X		
Village FE			X	
R-squared	0.48	0.38	0.54	0.24
Observations	1,413	1,413	1,413	1,413

The sample includes all farmers in the control group who completed all three meter reading survey rounds during the intervention. The outcome is monthly hours of irrigation by the farmer during the three intervention months (scaled to 31 days). Standard errors clustered at the randomization pair level are in brackets.

Table 13: Counterfactual Payments Under Alternative Benchmark Scenarios

	Average payment (INR per participant per month)
Actual payments, treatment group	1,238
Counterfactual payments, control group	
1. Benchmark = Pre-program value	1,872
2. Benchmark predicted from pre-program value	1,022
3. Benchmark predicted from basic characteristics	1,325
4. Benchmark predicted using rich characteristics	1,323
5. Benchmark set as in the treatment group	1,030